Google Cloud

Review: Fundamentals of Cloud Run

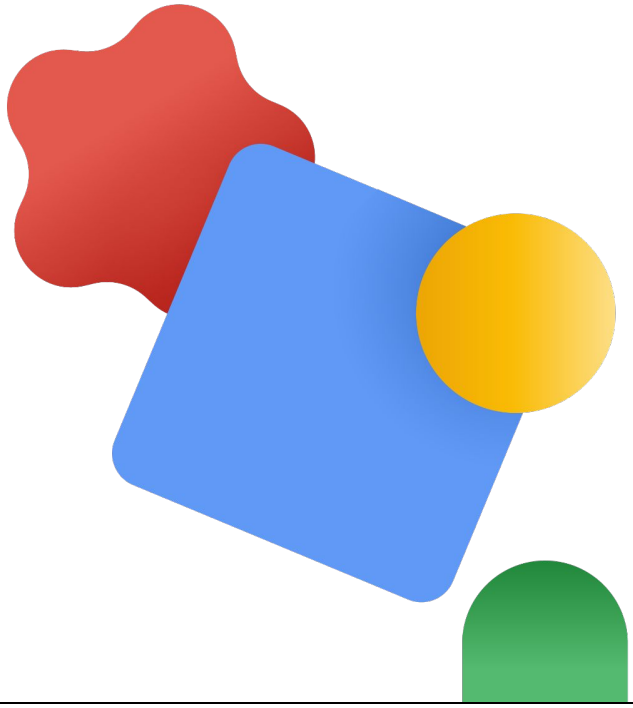

Let’s review the topics that were discussed in this module.

Proprietary + Confidential

In this module, you learned about:

01	Cloud Run overview
02	Resource model
03	Container lifecycle
04	Autoscaling
05	Access control

Google Cloud

We started this module by first introducing Cloud Run and provided an overview of how to use Cloud Run with services and jobs. We also discussed some of the features and benefits of using Cloud Run for your application workloads.

We then discussed the resource model that is used in Cloud Run for services and jobs. A service consists of service revisions, where each revision is scaled with multiple container instances to process requests that are sent to the service. A job consists of one or multiple independent tasks that are executed in parallel in a given job execution. Each task runs one container instance.

The various phases in the lifecycle of a container were discussed, followed by Cloud Run autoscaling. The module also discussed how to prevent scale to zero instances with minimum instance configuration, and discussed how to set configuration for maximum instances and maximum concurrency.

Finally, we discussed how to control access to your Cloud Run service or job with IAM, and with network ingress settings, and how to access internal resources with Serverless VPC Access.